One Day Hackathon: Intel Image Classification

TOC

Project objective

To create a convolutional neural network classifier that can differentiate between buildings, forests, oceans, mountains, streets, & glaciers

Overview

Data Collection

- 14,000 images in Train
- 3,000 images in Test*
- 7,000 images in Prediction**

Data Cleaning / EDA

- Dataset was clean, dealing only with images that were already sorted

Modeling and Interpretation

- Convolutional Neural Network
- Had most difficulty classifying Buildings against Streets, And Glaciers against Mountains
- 90%_ val_accuracy score

Our custom model performed at 74% accuracy

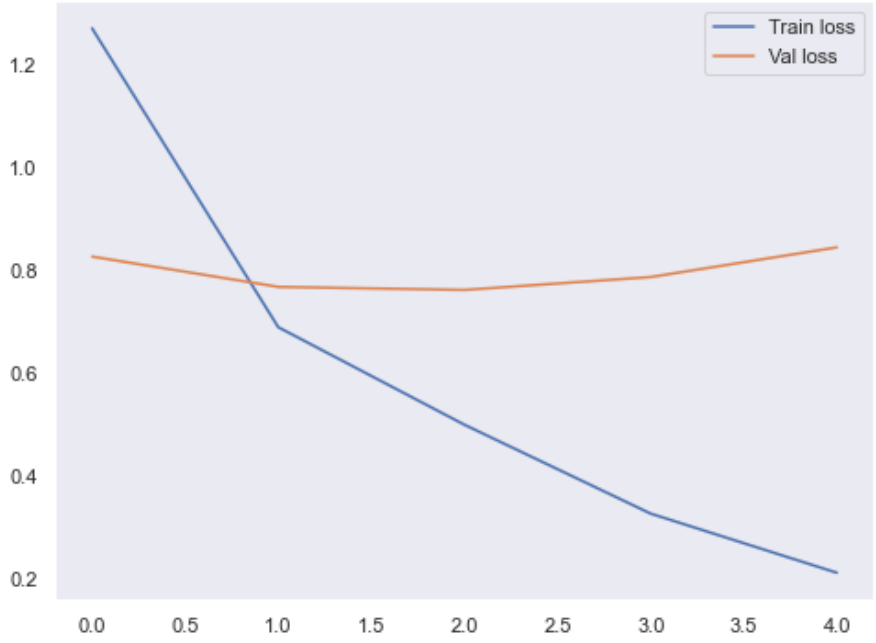

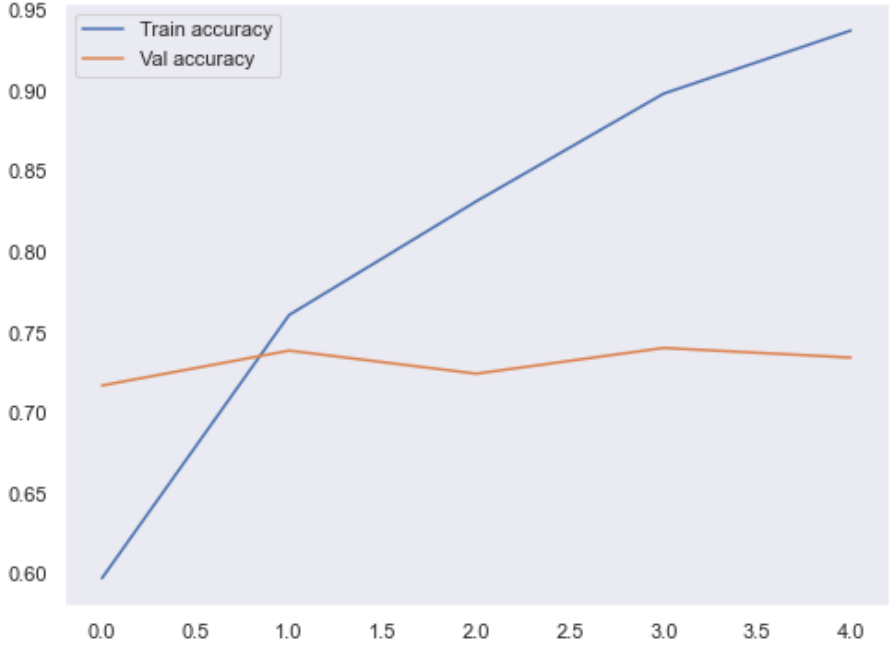

Baseline score for a 6 category classification problem with equal class representation is 16.6%

Transfer learning achieved 90%+ accuracy

Could have achieved marginally higher accuracy on test data with more training epochs, given more time

Model best identifies forests, worst at glaciers

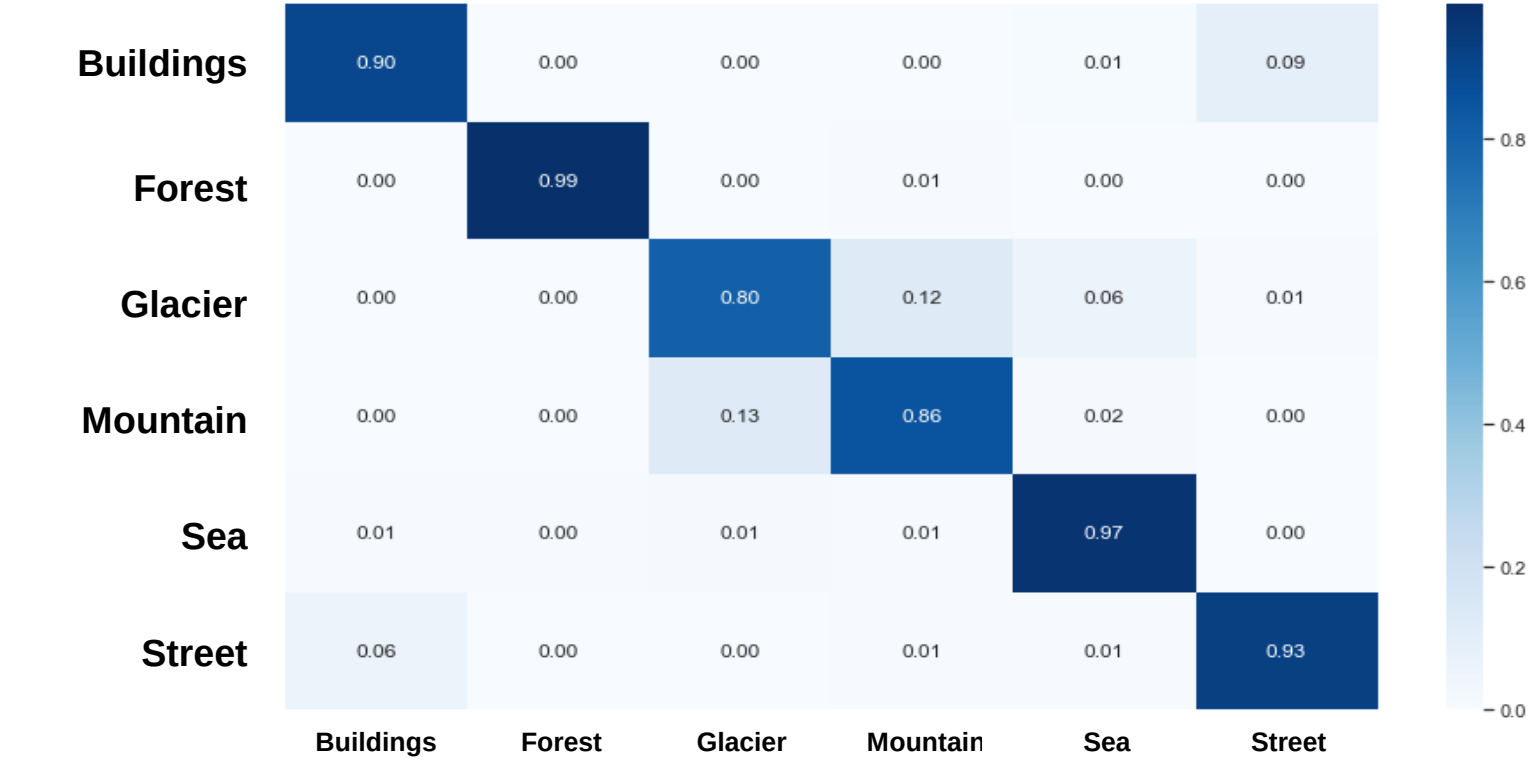

To emphasize how good 90% accuracy is on this task, let's play a game...

Let’s play a game: Which is which?

Mountain or glacier?

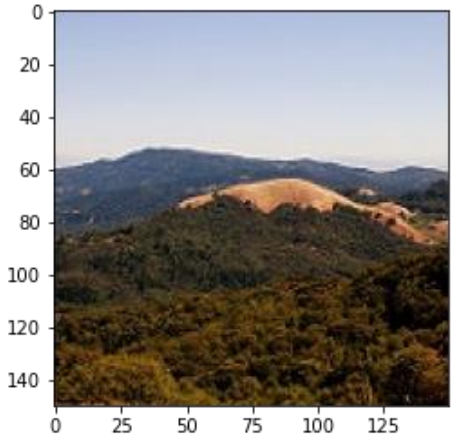

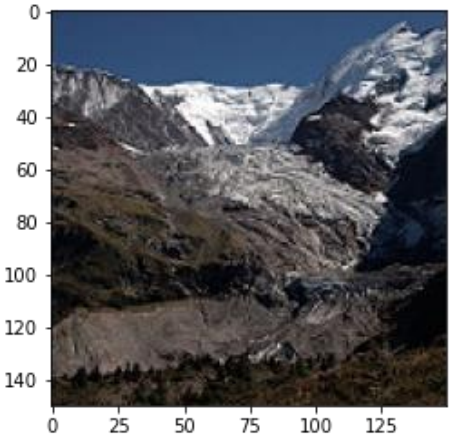

Left is mountain! And Right is glacier!

That was easy? Ok, how about these two?

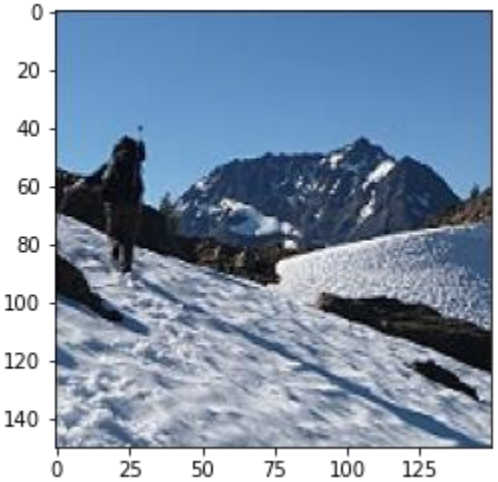

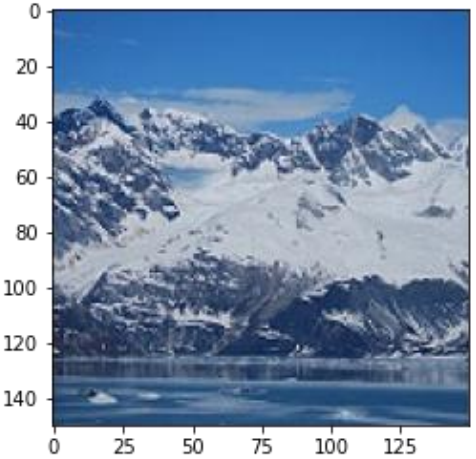

Left is mountain! And right is glacier

Let’s play a game: Which is which?

Building or street?

Street is on left Building is on right

Let’s play a game: Which is which?

Building or street?

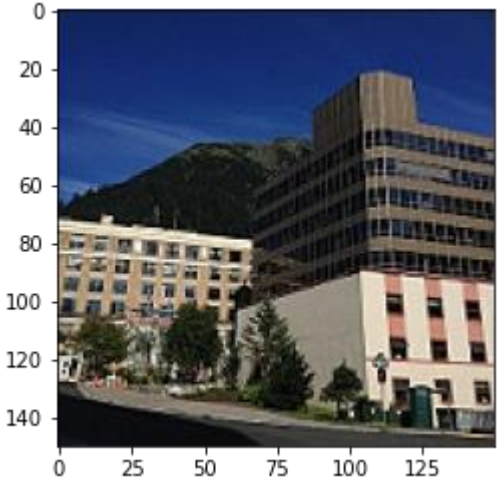

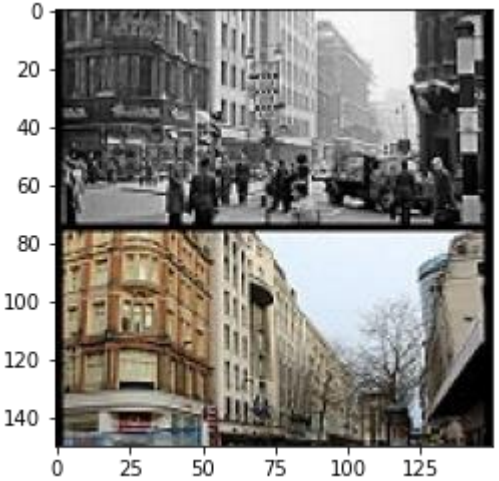

Building is on left and Street is on right!

Conclusion

Image classification is a powerful tool that requires complex architecture, extremely deep models with very large numbers of parameters, and huge datasets for effective training - but with transfer learning, you can look like a pro with just a few tweaks and custom layers!

Thank you.